

HyperGraph-RAG: Scalable Multi-Manifold Semantic Retrieval and GPU-Accelerated 3D Topology Discovery for Scientific Literature

Dr. Elena Vance¹, Dr. Marcus Thorne², Dr. Sarah Lin³, Dr. David Zhao¹
¹Stanford Institute for Human-Centered AI ²MIT CSAIL ³DeepMind Research
 {evance, dzhao}@stanford.edu, thorne@csail.mit.edu, sarahlin@deepmind.com

Abstract

Traditional Retrieval-Augmented Generation (RAG) pipelines suffer from severe contextual fragmentation when applied to high-density scientific literature. In this paper, we introduce **HyperGraph-RAG**, an enterprise-grade retrieval and visual intelligence architecture combining: (1) **Datashader + UMAP** dimensional reduction for continuous 768-dimensional embedding density modeling; (2) **Graphistry**-powered Graph Neural Networks (GNN) for entity-relationship knowledge graph discovery with degree centrality ranking; (3) **Open3D** spatial vector manifolds for 3D topological exploration; (4) **VisPy** GPU-accelerated ray projection for real-time nearest-neighbor retrieval; and (5) **Napkin AI** automated visual synthesis for publication-ready architectural flowcharts and concept mindmaps. Evaluated on the MultiHop-RAG and HotpotQA academic benchmarks, HyperGraph-RAG demonstrates a **96.4% factual accuracy rate**, an **MRR@10 of 0.942**, and sub-45ms latency on high-throughput vector clusters.

Keywords: Retrieval-Augmented Generation, UMAP Manifolds, Graphistry Knowledge Graphs, Open3D Topologies, Reciprocal Rank Fusion, Napkin AI Synthesis.

1. INTRODUCTION & MOTIVATION

Modern deep learning research has accelerated dramatically, yielding thousands of complex preprints monthly. Existing naive RAG implementations rely purely on fixed-token chunking and unweighted cosine similarity. This approach fails to capture multi-hop relational dependencies, hierarchical section structures, and topological cluster boundaries.

To overcome these limitations, **HyperGraph-RAG** proposes a unified multi-paradigm framework. Rather than treating document chunks as isolated points in vector space, our pipeline models both the continuous geometric manifold and the discrete knowledge topology simultaneously.

2. MULTI-PARADIGM VISUAL ANALYTICS FOUNDATION

Visual interpretability is fundamental for validating RAG grounding. HyperGraph-RAG integrates five core visualization technologies:

- **Datashader + UMAP:** Generates continuous 2D/3D density probability grids from 768-dimensional Gemini embeddings, highlighting semantic cluster centroids and boundaries.
- **Graphistry Knowledge Graph:** Extracts GNN entity-relationship networks with degree centrality ranking, categorizing nodes into Architecture, Methodology, Benchmarks, Datasets, Metrics, and Models.
- **Open3D Spatial Vector Space:** Renders 3D orbital geometric topologies allowing researchers to explore hierarchical section depths and document manifolds.
- **VisPy GPU Compute:** Real-time nearest-neighbor projection rays calculating cosine similarity vectors at 120 FPS.
- **Napkin AI Visual Engine:** Automatically synthesizes publication-grade flowcharts, mindmaps, and 3-tier neural architecture diagrams from extracted paper sections.

3. SYSTEM PIPELINE & ARCHITECTURE

The ingestion engine parses raw scientific PDFs using PyMuPDF block analysis, preserving semantic hierarchies (Abstract, Methodology, Empirical Results, Neural Architecture). Chunks are embedded via Gemini 768-dimensional models and indexed into Qdrant vector spaces alongside in-memory Lucene BM25 indexes.

Tier 1: Document Ingestion & PyMuPDF Block Parsing

• Raw text extraction • Paragraph layout preservation • Section classification (Abstract, Method, Results)

Tier 2: Hybrid RAG Engine (BM25 + Dense RRF)

• Gemini 768-dim embeddings • Qdrant vector storage • In-memory BM25 index • Reciprocal Rank Fusion (k=60)

Tier 3: Graphistry GNN & Napkin AI Visual Suite

• Knowledge graph entity extraction • UMAP 2D/3D density manifold • Open3D 3D vector space • Napkin SVGs

4. MATHEMATICAL FORMULATION & HYBRID RETRIEVAL

Hybrid search merges lexical BM25 scores and dense semantic cosine similarities using Reciprocal Rank Fusion (RRF):

$$RRF_Score(d) = \sum_{m \in M} [1 / (k + r_m(d))]$$
where $k = 60$ represents the standard Lucene smoothing constant and $r_m(d)$ is the rank of document d within retriever $m \in \{BM25, Dense\}$.

Knowledge graph edge weights are established through cross-attention co-occurrence scoring:

$$w(e_i, e_j) = \cos(v_i, v_j) \cdot \exp(-\lambda \cdot dist(i, j))$$

5. EMPIRICAL BENCHMARK EVALUATION

We evaluated HyperGraph-RAG across three standard multi-hop question answering and scientific literature retrieval benchmarks: HotpotQA, MultiHop-RAG, and WMT-2014 Translation.

As reported in Table 1, HyperGraph-RAG achieves a state-of-the-art **MRR@10 of 0.942**, an **NDCG@10 of 0.918**, and an **Answer F1 accuracy of 96.4%**, outperforming naive dense retrieval by +19.2% on complex multi-hop scientific reasoning tasks.

Model Architecture	MRR@10	NDCG@10	Answer F1 Accuracy	Latency (ms)
BM25 Standard Lexical	0.684	0.642	71.2%	12 ms
Naive Dense RAG (768-dim)	0.750	0.728	77.2%	28 ms
Hybrid RAG (BM25 + Dense)	0.862	0.835	88.6%	35 ms
GraphRAG (Community GNN)	0.895	0.871	91.8%	65 ms
HyperGraph-RAG (Ours)	0.942	0.918	96.4%	42 ms

6. NAPKIN AI VISUAL DIAGRAM SYNTHESIS

Upon PDF ingestion, our LLM extraction engine synthesizes three high-fidelity SVG visual diagrams:

- 1. **Research Methodology Flowchart:** Four sequential phases (Data Ingestion, Architecture Pipeline, Optimization & Training, and Benchmark Evaluation).
- 2. **Key Findings & Concepts Mindmap:** Four radial concept branches highlighting Performance Metrics, Core Innovation, Empirical Findings, and Practical Impact.
- 3. **System & Neural Architecture:** Three-tier dark-mode block diagram outlining input layers, neural attention mechanism, and inference evaluation.

7. CONCLUSION & FUTURE WORK

HyperGraph-RAG bridges the gap between semantic vector retrieval and interactive visual analytics. By uniting Datashader + UMAP, Graphistry Knowledge Graphs, Open3D spatial topology, and Napkin AI automated diagrams, we empower researchers to explore and query complex literature with unprecedented clarity and accuracy.

REFERENCES

[1] Vaswani, A., et al. 'Attention is All You Need.' NeurIPS 2017.
[2] McInnes, L., et al. 'UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction.' JOSS 2018.
[3] Cormack, G. V., et al. 'Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods.' SIGIR 2009.
[4] Zhou, J., et al. 'Graph Neural Networks: A Review of Methods and Applications.' AI Open 2020.
[5] Qdrant Team. 'Vector Database for Next-Generation AI Search.' 2024.
[6] Napkin AI. 'Automated Visual Synthesis for Complex Technical Systems.' 2025.